

EEC 281 - Homework/Project #4

Work individually, but I strongly recommend working with someone in the class nearby so you can help each other when you get stuck, with consideration of the [Course Collaboration Policy](#). Send an email to me if something is not clear and I will update the assignment using **green font**.

- **Submit:** (1) all code you wrote (not generated or provided files) including verilog hardware, verilog testing, matlab, etc. (2) other requested items such as diagrams etc.
 - i. Upload a single pdf to <https://canvas.ucdavis.edu/>.
 - ii. Place all of your answers and code into a single pdf file with all problems and material *in order* (i.e., problem 1, problem 2,...).
 - iii. Add titles to pages and file names so it is clear to which problem they belong. For example, Problem 1, prob1.v, prob1.vt,...
- **Diagrams.** If a problem requires a diagram, include details such as datapath, memory, control, I/O, pipeline stages, word widths in bits, etc. There must be enough detail so that the exact *functional* operation of the block can be determined by someone with a reasonable knowledge of what simple blocks do. A satisfactory diagram may require an extra-large sheet.
- **Verilog.** If a problem requires a verilog design, turn in copies of both hardware and test verilog code.
 - *** Where three '*'s appear in the description, perform the required test(s) and turn in a printout of either:
 1. a table printed by your verilog testbench module listing all inputs and corresponding outputs,
 2. a simvision waveform plot which shows (labeled and highlighted) corresponding inputs and outputs, or
 3. verilog test code which compares a) your hardware circuit and b) a simple [Golden Reference](#) circuit (using high-level functions such as "+"). Include two copy & paste sections of text from your simulation's output (one section showing a large number of passes, and one small section showing where you purposely make a *very small* change to either your designed hardware circuit or your reference circuit to force the comparison to fail). It should look something like this:

```
input=0101, out_hw=11110000, out_ref=11110000, ok
input=0111, out_hw=11110001, out_ref=11110001, ok
...

input=0101, out_hw=11110000, out_ref=11110001, Error!
...
```
 - For 1 and 3, the output must be copied & pasted directly from the simulator's output without any modifications.
 - In all cases, **Show how you verified** the correctness of your simulation's outputs.
 - Keep "hardware" modules separate from testing code. Instantiate a copy of your processing module(s) in your testing module (the highest level module) and drive the inputs and check the outputs from there.
 - Your verilog must implement *hardware* and be cleanly synthesizable, and follow guidelines in the verilog handouts. For example, having a `for` loop in your "hardware" verilog will result in an automatic **80%** reduction in points since the verilog is therefore not implementing hardware.
- **Synthesis.** If a problem requires synthesis, turn in copies of the following. Print in a way that results are easy to understand. Delete sections of many repeated lines with a few copies of the line plus the comment: `<many lines removed> .`
 1. `dc-<filename>.tcl` (or equivalent)
 2. `*.area` file
 3. `*.logs` file (not `command.log`) Edit and reduce "Beginning Delay Optimization Phase" and "Beginning Area-Recovery Phase" sections.
 4. `*.tim` file; first (longest) path only

The `"always @(*)"` verilog construct may be used.

Run all compiles with "medium" effort. Do not modify the synthesis script except for functional purposes (e.g., to specify source file names).

- **Functionality.** For each design problem, clearly state: 1) whether the design is fully functional, and 2) the failing sections if any exist.
- **Point deductions/additions.** `TotalProbPts` is the sum of all points possible.
 - [Up to `TotalProbPts × 50%`] point reduction for not plainly certifying/showing that your circuit is **functionally correct**. This sounds drastic but you should have checked the correctness of your circuits' outputs anyway, it is impractical for the grader to check every result of every submission by eye, and thus an un-certified design will be treated like a marginally-functional design after a cursory glance at the hardware. In the worst case if there is no indication a design works at all, zero points will be given.

Following is an example of a fine way to certify correctness, if the "Y/N" is written either a) by hand individually for each test or b) automatically with a golden reference checker; but not printed automatically without individual checking.

inA	inB	outExp	outMantissa	I Certify Correct
10101100	00110101	110010	01100110100101	Y
00000101	10110101	101010	01010101010101	Y
01010100	11101010	010100	11010101100101	no // this indicates I recognize there is an error here

- [Up to `TotalProbPts × 10%`] point reduction if parts of different problems are mixed up together (please don't do it; it makes grading much more difficult than you probably realize)
- [Up to `TotalProbPts × 10%`] extra credit will be given for especially thorough, well-documented, or insightful solutions.
- **Clarity.** For full credit, your submission must be easily understandable and well commented. Print code and CAD reports using a mono-spaced font (e.g., *Courier*) and a small size such as 9 point.

Total: 375 points

1. [90 pts] This problem requires the design of a block which calculates $\tan(\theta)$ for a given θ , every clock cycle. Theta ranges from 0 to slightly less than $\pi/4$, or 45° . The latency may be as many cycles as needed.

The block's I/O signals are:

- `theta` input
12-bit fixed-point unsigned where:
0000_0000_0000 = 0.000 radians, and
1111_1111_1111 = $\pi/4 \cdot (4095/4096)$ radians
- `out` output
is 16-bit fixed-point 2's complement.

The test procedure is as follows:

1. Generate all 2^{12} possible theta inputs from your verilog testbench
2. Calculate the output for each theta with your verilog design
3. Output both the a) theta input and b) verilog output to a matlab-readable *.m file
4. Compare a) the verilog output and b) the matlab expression `tan(theta)` using `diff.m`

Design and write verilog for the block two ways:

1. With a "full" lookup table (for an area reference point)
2. Same as the previous design but where outputs for odd theta inputs are linearly-interpolated between samples from the lookup table(s) and the system has a throughput of one calculation for every two clock cycles. The goal is smaller area.

For each of your two designs, submit (a) through (d) below. When submitting the verilog file of your large lookup table, print only the first ~25 lines and the last ~25 lines and insert the comment "<Many lines removed>" for lines you deleted.

- a) [10 pts] Detailed block diagram with all functional details.
- b) [10 pts] Pipelined block diagram.

c) [10–25 pts] (Accuracy points for smallest error compared to matlab, in comparison to other working designs in the class.)

Write the $\text{Energy_diff}/\text{Energy_data0}$ value in dB in your report and also submit the four plots produced by `diff.m`.


d) [45 pts] Synthesize your design at the following three [cycle time values](#) and report the 1) *achieved* cycle time (and corresponding clock frequency) and 2) area for each:

1. a very long cycle time, e.g., 1 ms (1 KHz), to find the minimum area;
2. a very short cycle time, e.g., 0.1 ns (10 GHz), to find the minimum cycle time;
3. a synthesis run with the cycle time set to the result from part (d)(2) multiplied times 1.5


Points for (c) and (d) are possible only if the design is fully functional and without synthesis errors or serious warnings.

[285 pts] The Alexnet convolutional neural net is widely credited with the dramatic rise in popularity of neural nets. Read the [2012 Alexnet paper](#) paying particular attention to Sections 1, 2, 3.4, and 3.5. This project consists of building and synthesizing custom hardware for the first convolutional layer of Alexnet using a reduced image size. The primary specifications for this simplified project are as follows:

- input image size of 35 pixels x 35 pixels x 1 grayscale "color". Each pixel is an 8-bit unsigned integer.
- filter size of 11 pixels x 11 pixels. Each filter coefficient is an 8-bit 2's complement integer.
- 49 2-dimensional 11 x 11 convolutions total with a stride of 4 pixels in each direction for each convolution. The first 3 convolutions of the first two rows are shown in the figure below.
- the output of each convolution is passed through a ReLU rectified linear function $\max(0, x)$.
- a Max Pooling ([article showing non-overlapped pooling](#)) calculation with $s = 2$ and $z = 3$ (overlapped) is performed on the output of the ReLU calculations for which you have all the necessary data (9 outputs).



First Row of Convolutions



Second Row of Convolutions

The testing environment is built as follows:

- the testbench performs the following steps in order one after another:
 1. it generates random data for the pixel memory and inputs one pixel (8 bits) at a time into the processor,
 2. it generates random data for the filter one pixel coefficient (8 bits) at a time into the processor, and
 3. starts the processor with an input *go* signal which starts a hardware controller located entirely within the processor.

For steps #1 and #2 which load data, there must be some coordination between the hardware and testbench however calculations made in the testbench must be minimized, for example the testbench could provide only a *valid* signal along with the data or some other similarly-simple interface. Control signals for the processing hardware described in step #3 must be located entirely within the hardware—none can be located in the testbench.

- For the first few samples of both the pixels and filter, use multiple combinations of 0 and +1 (for both **pixels and filter coeffs**), -1, -128 and +127 (for filter coeffs), and +255 (for pixels), to check calculations quickly.
- See the "verilog: example code" web page for example code. Use random seed "123".
- all pixel data, filter data, and outputs are printed to a *.m file for analysis in matlab.

Other requirements:

- all convolutions should be calculated in approximately **539** clock cycles.
- ReLU and max pooling calculations must be overlapped with convolutional calculations as much as is easily achievable
- use only single-ported memories (modeled by verilog memories with a single address bus for both reads and writes)
- You may use "signed" wires and regs for signed data calculations
- use "*" for multipliers, and "+" for all adders.
- Because synthesis times will be too long (approx 30 minutes) if all memories are synthesized, place all memory modules outside the top-level processor, registering values immediately before leaving and after entering the main module.
- Minimum pipeline stages: convolution processor (2 stages), ReLU (1 stage), Max pooling (1 stage)

Submit the following.

- a) [25 pts] A bulleted list and a few sentences describing your design including such features as: the number and size of memories, number of multipliers and adders, exact number of cycles to complete all convolutions, number of pipeline stages, in what order are the pieces and in what order are the convolutions calculated, etc.
- b) [25 pts] Detailed pipelined block diagram with all functional details.
- c) [25 pts] Detailed timing diagram showing the timing of the various processing units
- d) [75 pts] Results for all convolutions printed by your verilog testbench.
- e) [50 pts] Matlab copy and pasted output showing whether your design matches the matlab model or not. Your hardware must exactly match the output of the [alex35.m matlab model](#).
- f) [75 pts] Synthesize your design at the following three [cycle time values](#) and report the 1) *achieved* cycle time (and corresponding clock frequency) and 2) area for each:
 1. a very long cycle time, e.g., 1 ms (1 KHz), to find the minimum area;
 2. a very short cycle time, e.g., 0.1 ns (10 GHz), to find the minimum cycle time;
 3. a synthesis run with the cycle time set to the result from part (f)(2) multiplied times 1.5
- g) [10 pts] Report the number of clock cycles between the *go* signal and 1) the first output, and 2) the final output.

Points for (e), (f), and (g) are possible only if the design is fully functional and without synthesis errors or serious warnings.

Hint: See the Synthesis handout for details on the *achievable cycle time* and reading synthesis timing reports.

Hint: See the "matlab: tips for 281" web page for suggestions on addressing memories in matlab.

2025/03/04 Posted
2025/03/05 Added max pooling reference link
2025/03/10 Clarified test pixel and filter samples
2025/03/12 Changed image size from 31x31 to 35x35 to give 9 outputs instead of 4.
2025/03/18 Fixed typo, removed "fourth" from problem 2(g)(2)